SYNTHETIC SAMPLE PAPER - fictional, for pipeline testing only

LoRA Adapters for Domain Adaptation in Machine Translation

Abstract
We study domain adaptation and low-resource machine translation for Amharic using LoRA adapters. Experiments use the FLORES-200 dataset.

1 Introduction
Full fine-tuning of large models is expensive for small teams. Low-rank adapters update only a small number of parameters.

2 Method
We insert LoRA adapters into a pretrained translation model and train only the adapters. The adapters transfer to the medical domain with small additional training.

3 Results
LoRA reaches 19.8 BLEU on FLORES-200 for Amharic.

4 Limitations
We evaluate only one language pair, Amharic to English.

5 Future Work
We plan to extend the study to dialectal data such as Algerian Arabic.